

NETWORKING FUNDAMENTALS

George Porter
Jan 9, 2019

UC San Diego

1



ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license
- These slides incorporate material from:
 - Alex C. Snoeren, UC San Diego
 - Michael Freedman and Kyle Jamieson, Princeton University
 - Internet Society
 - Computer Networking: A Top Down Approach 6th edition

UC San Diego

2

ANNOUNCEMENTS

Peterson & Davie 1.3 (network architectures), 1.4, 1.5, 2.5, 3.2

Class Q&A forum (via Google) now available (link off course web page)

UC San Diego

3

Outline

1. Packet switching
2. Addressing
3. Performance
4. TCP and sockets
5. Reliable transmission



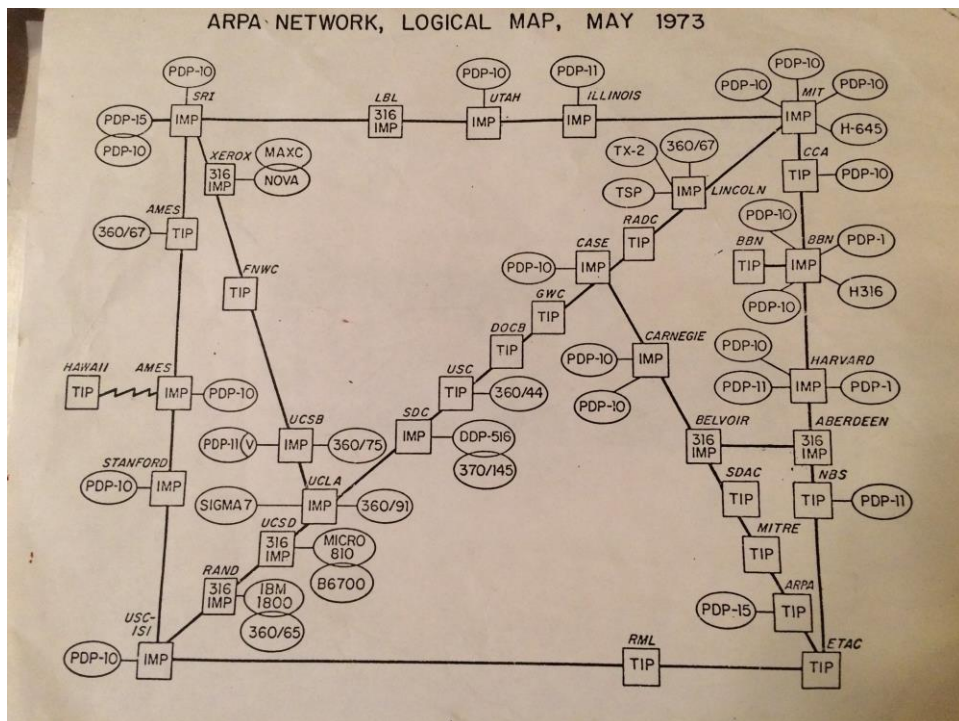
4

BRIEF HISTORY OF THE INTERNET

- 1968 - DARPA (Defense Advanced Research Projects Agency) contracts with BBN (Bolt, Beranek & Newman) to create ARPAnet
- 1970 - First five nodes:
 - UCLA
 - Stanford
 - UC Santa Barbara
 - U of Utah, and
 - BBN
- 1974 - TCP specification by Vint Cerf
- 1984 – On January 1, the Internet with its 1000 hosts converts en masse to using TCP/IP for its messaging

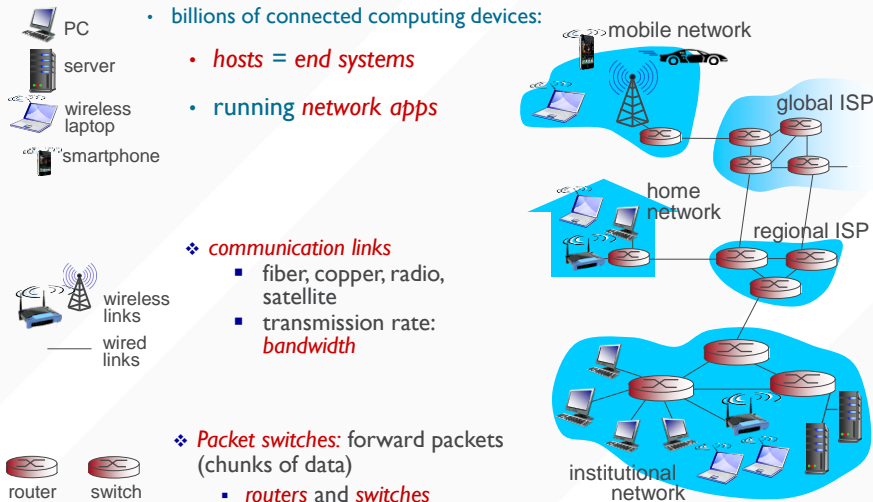
Data from the Internet Society

5



6

AN INTER-NETWORK




7

PACKET SWITCHING


- Data transmitted from source to destination in discrete *packets*
 - Variable size, usually a maximum transmission unit (MTU) of ~1500 bytes
 - 1 GB file is approx 715,000 packets
- Each packet routed independently
 - Has its own source and destination *address* used to find the route to the ultimate destination

8

NETWORK TERMINOLOGY




A single packet-switched network



Interconnection of networks

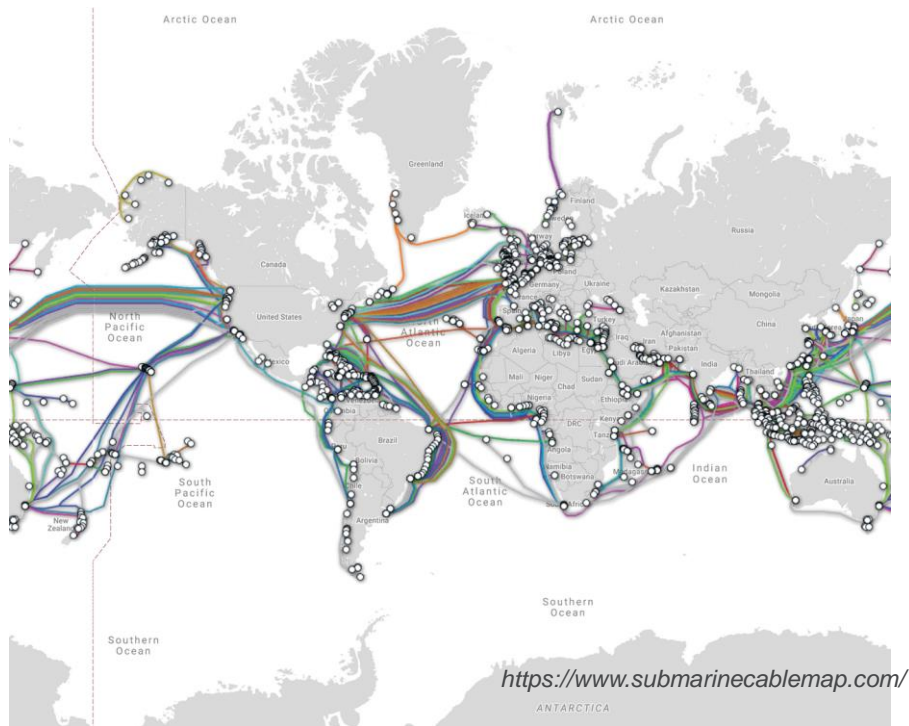
9

ROUTING PACKETS BETWEEN NETWORKS



- Networks use *Border Gateway Protocol (BGP)* to announce reachability
- Each network talks just with its neighbors
- Goal is to get a packet to the destination network
- It is up to that destination network to get individual packets to their ultimate destination
- Back-to-back packets from the same "connection" might take different paths!
- Might arrive out of order too


10



11

ROUTING PACKETS WITHIN A SINGLE NETWORK


- E.g. UCSD, San Diego Comcast, a Google datacenter
- Packet switches communicate with each other using an *intra-domain* protocol called a Link-state protocol
 - Each packet switch sends a list of its neighbors to every other switch via a *broadcast*
 - As a result, each switch has the same “map” of what the network’s *topology* looks like



12

ROUTING PACKETS WITHIN A SINGLE NETWORK


- Each packet switch uses its local topology map to choose paths to each destination
- Typically using Dijkstra's shortest path algorithm
- Packets are forwarded along this shortest path *tree*




13

RESOURCE SHARING

Multiplexing multiple logical flows over a single physical link



- Resource: links and nodes
- How to share a link?
 - Multiplexing
 - De-multiplexing
 - Queueing



14

Outline

1. Packet switching
2. Addressing
3. Performance
4. TCP and sockets
5. Reliable transmission



15

ADDRESSING CONSIDERATIONS

- Fixed length or variable length addresses?
- Issues:
 - Flexibility
 - Processing costs
 - Header size
- Engineering choice: IP uses fixed length addresses
- 32-bits in an IPv4 address
 - Dotted decimal format a.b.c.d
 - Each represent 8 bits of address
- Hierarchical: Network part and host part
 - E.g. IP address 128.54.70.238
 - 128.54 refers to the UCSD campus network
 - 70.238 refers to the host ieng6.ucsd.edu
- Which part is network vs. host?

16

CLASS-BASED ADDRESSING

- Most significant bits determines “class” of address

Class A

0	Network	Host
---	---------	------

127 nets, 16M hosts

Class B

		14	16
1	0	Network	Host

16K nets, 64K hosts

Class C

21			8	
1	1	0	Network	Host

2M nets, 254 hosts

- Special addresses
 - Class D (1110) for multicast, Class E (1111) experimental
 - 127.0.0.1: local host (a.k.a. the loopback address)
 - Host bits all set to 0: network address
 - Host bits all set to 1: broadcast address

17


IP FORWARDING TABLES

- Router needs to know where to forward a packet
- Forwarding table contains:
 - List of network names and **next hop** routers
 - Local networks have entries specifying which interface
 - Link-local hosts can be delivered with Layer-2 forwarding
- E.g. www.ucsd.edu address is 132.239.180.101
 - Class B address – class + network is 132.239
 - Lookup 132.239 in forwarding table
 - Prefix – part of address that really matters for routing

18

SUBNETTING

- Individual networks may be composed of several LANs
 - Only want traffic destined to local hosts on physical network
 - Routers need a way to know which hosts on which LAN
- Networks can be arbitrarily decomposed into **subnets**
 - Each subnet is simply a prefix of the host address portion
 - Subnet prefix can be of any length, specified with **netmask**



19

SUBNET ADDRESSES


- Every (sub)network has an address and a netmask
 - Netmask tells which bits of the network address is important
 - Convention suggests it be a proper prefix
- Netmask written as an all-ones IP address
 - E.g., Class B netmask is 255.255.0.0
 - Sometimes expressed in terms of number of 1s, e.g., /16
- Need to size subnet appropriately for each LAN
 - Only have remaining bits to specify host addresses

20

IP ADDRESS PROBLEM (1991)

- Address space depletion
 - In danger of running out of classes A and B
- Why?
 - Class C too small for most organizations (only ~250 addresses)
 - Very few class A – very careful about giving them out (who has 16M hosts anyway?)
 - Class B – greatest problem

21



22

CIDR

- Classless Inter-Domain Routing (1993)
 - Networks described by variable-length prefix and length
 - Allows arbitrary allocation between network and host address




- e.g. 10.95.1.2 contained within 10.0.0.0/8:
 - 10.0.0.0 is network and remainder (95.1.2) is host
- Pro: Finer grained allocation; aggregation
- Con: More expensive lookup: **longest prefix match**

23

FORWARDING TABLE EXAMPLE (R2)

- Packet to 10.1.1.6
- Matches 10.1.0.0/23

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.16.1
10.1.8.0/24	interface1
10.1.2.0/23	interface2
10.1.0.0/23	10.1.2.2
10.1.16.0/24	interface3




24

FORWARDING TABLE EXAMPLE 2

- Packet to 10.1.1.6
- Matches 10.1.1.4/30
 - Longest prefix match

Routing table at R1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.2.1
10.1.0.0/24	interface1
10.1.1.0/24	interface2
10.1.2.0/23	interface3
10.1.1.4/30	10.1.1.101




25

FORWARDING TABLE EXAMPLE 3

- Packet to 10.1.1.6
- Direct route
 - Longest prefix match

Routing table at H1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.1.1
10.1.1.0/24	interface1
10.1.1.4/30	interface2



26


ASSIGNING ADDRESSES VIA DHCP

- DHCP server is responsible for providing configuration information to hosts
- There is at least one DHCP server for an administrative domain
- DHCP server maintains a pool of available addresses

27

DHCP IN ACTION

- Newly booted or attached host sends DHCPDISCOVER message to a special IP address (255.255.255.255)
- DHCP relay agent unicasts the message to DHCP server and waits for the response



28

DNS HOSTNAME VERSUS IP ADDRESS

- **DNS host name** (e.g. www.cs.ucsd.edu)
 - **Mnemonic** name appreciated by humans
 - **Variable length**, full alphabet of characters
 - Provides **little** (if any) information about **location**
- **IP address** (e.g. 128.112.136.35)
 - Numerical address appreciated by **routers**
 - **Fixed length**, decimal number
 - **Hierarchical** address space, related to host **location**

29

MAPPING NAMES TO ADDRESSES

```

GETADDRINFO(3)                Linux Programmer's Manual                GETADDRINFO(3)

NAME
    getaddrinfo, freeaddrinfo, gai_strerror - network address and service
    translation

SYNOPSIS
    #include <sys/types.h>
    #include <sys/socket.h>
    #include <netdb.h>

    int getaddrinfo(const char *node, const char *service,
                    const struct addrinfo *hints,
                    struct addrinfo **res);

    void freeaddrinfo(struct addrinfo *res);

    const char *gai_strerror(int errcode);
  
```

30

LINKED LIST OF 'ADDRINFO' STRUCTS

```
struct addrinfo {
    int             ai_flags;
    int             ai_family;
    int             ai_socktype;
    int             ai_protocol;
    socklen_t       ai_addrlen;
    struct sockaddr *ai_addr;
    char            *ai_canonname;
    struct addrinfo *ai_next;
};
```

- Q: Why a linked list?
- Q: Which of the multiple results should you use?

31


Outline

1. Packet switching
2. Addressing
3. Performance
4. TCP and sockets
5. Reliable transmission



32

PERFORMANCE METRICS



- Bandwidth: number of bits transmitted per unit of time
- Latency = Propagation + Transmit + Queue
 - Propagation = Distance/SpeedOfLight(*)
 - Transmit = 1 bit/Bandwidth
- Overhead
 - # secs for CPU to put message on wire
- Error rate
 - Probability P that message will not arrive intact

* In that particular medium

33

BANDWIDTH VS. LATENCY

1 Byte Object

	Latency: 1 ms	Latency: 100 ms
Bandwidth: 1 Mbps	1,008 μ s	100,008 μ s
Bandwidth: 100 Mbps	1,000 μ s	100,000 μ s

10 MB Object

	Latency: 1 ms	Latency: 100 ms
Bandwidth: 1 Mbps	80.001 s	80.1 s
Bandwidth: 100 Mbps	.801 s	.9 s

34

TERMINOLOGY STYLE

- Mega versus Mega, Kilo versus Kilo
 - Computer architecture: Mega $\rightarrow 2^{20}$, Kilo $\rightarrow 2^{10}$
 - Computer networks: Mega $\rightarrow 10^6$, Kilo $\rightarrow 10^3$
- Mbps versus MBps
 - Networks: typically megabits per second
 - Architecture: typically megabytes per second
- Bandwidth versus throughput
 - Bandwidth: available over link
 - Throughput: available to application

35

Outline

1. Packet switching
2. Addressing
3. Performance
4. TCP and sockets
5. Reliable transmission



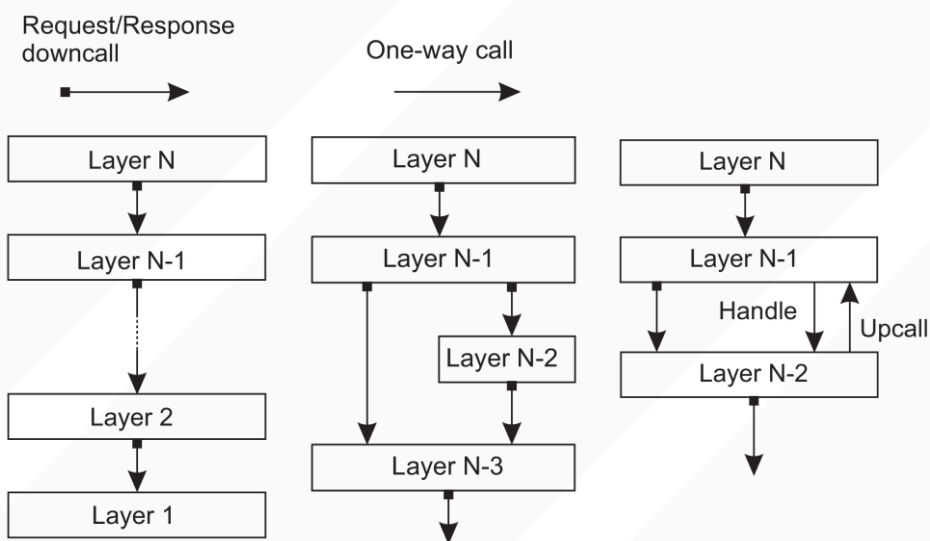
36

LAYERING: A MODULAR APPROACH TO COMM

- Sub-divide responsibilities
 - Each layer relies on services from layer below
 - Each layer exports services to layer above
- Interface between layers defines interaction
 - Hides implementation details (encapsulation)
 - Layers can change without disturbing other layers (modularity)
- Interface among peers in a layer is a **protocol**
 - If peers speak same protocol, they can interoperate


37

LAYERED ARCHITECTURES



38

SERVICE AND PROTOCOL INTERFACES



39

WHAT ARE PROTOCOLS?



Image from public domain


- Explicit and implicit conventions for how to communicate
 - Not for what is communicated



40

WHERE DO PROTOCOLS COME FROM?

- Standards bodies
 - IETF: Internet Engineering Task Force
 - ISO: International Standards Organization
- Community efforts
 - "Request for comments"
 - Bitcoin
- Corporations/industry
 - RealAudio™, Call of Duty multiplayer, Skype



41

HOW ARE PROTOCOLS SPECIFIED?

Prose/BNF

3.2. HEADER FIELD DEFINITIONS


These rules show a field meta-syntax, without regard for the particular type or internal syntax. Their purpose is to permit detection of fields; also, they present to higher-level parsers an image of each field as fitting on one line.

```


field      = field-name ":" [ field-body ] CRLF
field-name = 1*<any CHAR, excluding CTLs, SPACE, and ">
field-body = field-body-contents
              [CRLF LWSP-char field-body]

field-body-contents =
  <the ASCII characters making up the field-body, as
  defined in the following sections, and consisting
  of combinations of atom, quoted-string, and
  specials tokens, or else consisting of texts>
  
```

State transition diagrams



Message Sequence Diagram



By Stefan Birkner, cc-by-sa-2.5,2.0,1.0

Packet formats

0	1	2	3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1	1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1	2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1	3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1

Version IHL Type of Service Total Length			

Identification Flags Fragment Offset			

Time to Live Protocol Header Checksum			


Source Address			

Destination Address			

Options Padding			

42


ROLE OF LAYERING IN PROTOCOLS



- Each layer offers useful semantics to layer above
 - IP gets packets to a destination host/server on the Internet (but is unreliable)
 - TCP uses IP to offer *reliable, in-order bytestream* abstraction
 - TCP useful for file transfer, as well as HTTP/web

43

INTERNET PROTOCOL SUITE



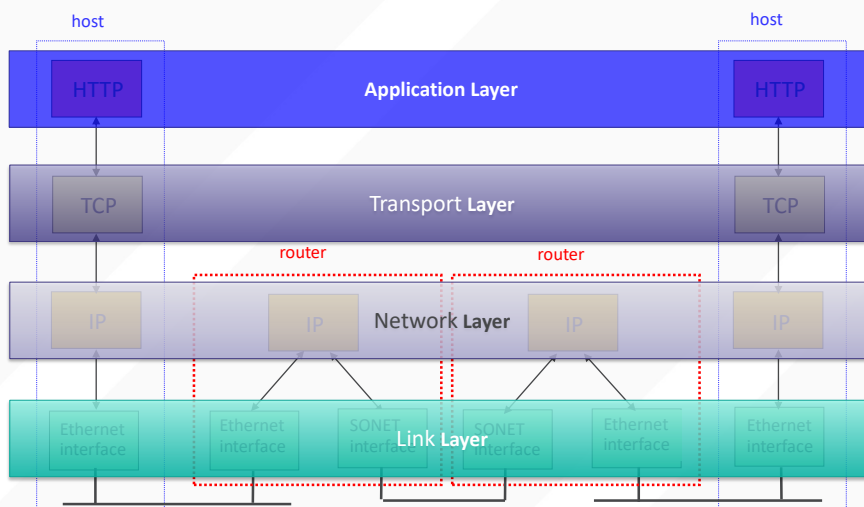
44

TRANSMISSION CONTROL PROTOCOL (TCP)

- Remember 1 GB \sim 715,000 packets?
 - Don't want to keep track of each packet, whether it got there, did it get lost? Did some get reordered??
- TCP offers *infinite bytestream* abstraction
 - If you put N bytes into TCP connection as sender, those N bytes will arrive to the destination in order, without loss, and without corruption
 - Compelling abstraction for higher-level applications such as web, video, gaming, ...

45

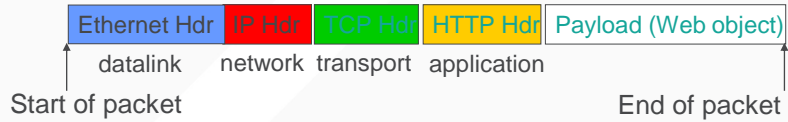
TCP/IP PROTOCOL STACK



46

ENCAPSULATION VIA HEADERS

- Typical web packet:



- Notice that layers add overhead
 - Space (headers), effective bandwidth
 - Time (processing headers), latency

47

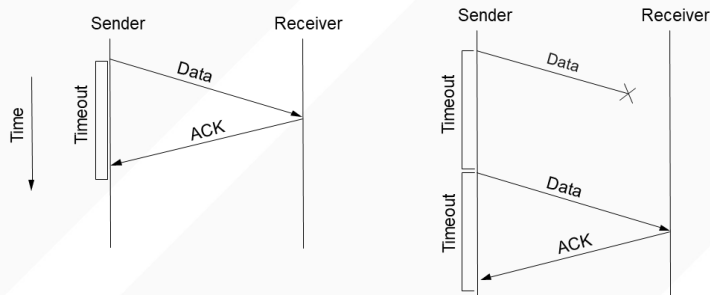
Outline

1. Packet switching
2. Addressing
3. Performance
4. TCP and sockets
5. Reliable transmission



48

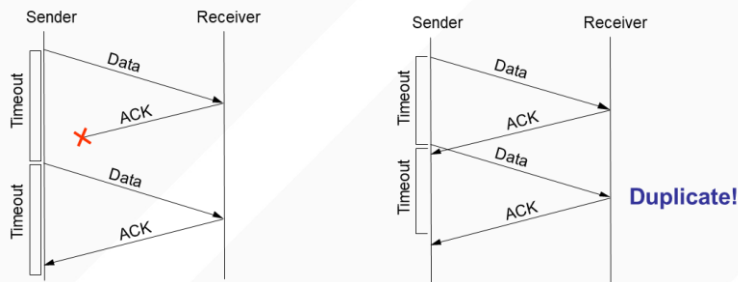
SIMPLE IDEA: ARQ



- Receiver sends acknowledgments (ACKs)
 - Sender “times out” and retransmits if it doesn’t receive them
- Basic approach is generically referred to as Automatic Repeat Request (ARQ)

49

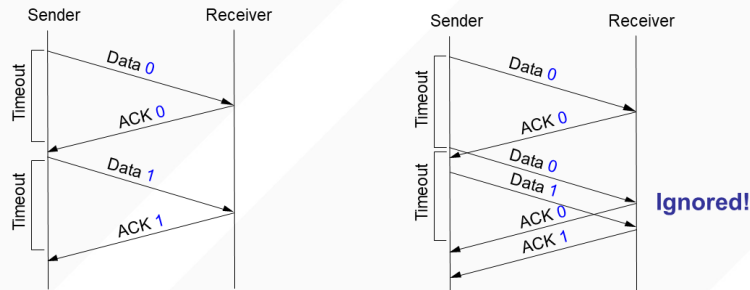
NOT SO FAST!



- Loss can occur on ACK channel as well
 - Sender cannot distinguish data loss from ACK loss
 - Sender will retransmit the data frame
- ACK loss—or early timeout—results in duplication
 - The receiver thinks the retransmission is new data

50

SEQUENCE NUMBERS



- Sequence numbers solve this problem
 - Receiver can simply ignore duplicate data
 - But must still send an ACK! (Why?)
- Simplest ARQ: **Stop-and-wait**
 - Only one outstanding frame at a time

51

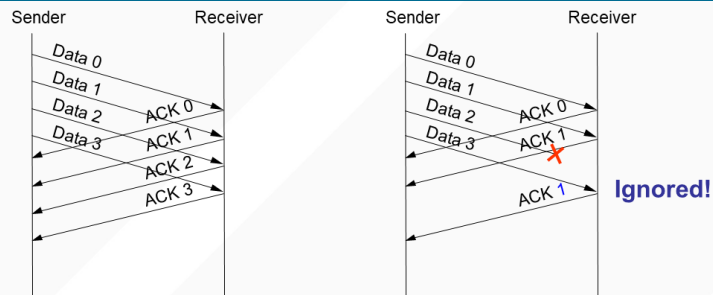
STOP AND WAIT PERFORMANCE



- Lousy performance if $x_{\text{mit}} 1 \text{ pkt} \ll \text{prop. delay}$
 - How bad?
- Want to utilize all available bandwidth
 - Need to keep more data "in flight"
 - How much? Called the **bandwidth-delay product**
- Also limited by quality of timeout (how long?)

52


PIPELINED TRANSMISSION



- Keep multiple packets “in flight”
 - Allows sender to make efficient use of the link
 - Sequence numbers ensure receiver can distinguish frames
- Sender buffers outstanding un-acked packets
 - Receiver ACKs the highest *consecutive* frame received
 - ACKs are *cumulative* (covers current frame and all previous)

53


SLIDING WINDOW: SENDER



- Window bounds outstanding unACKed data
 - Implies need for buffering at sender
- “Last” ACK applies to *in-order* data
- What to do on a timeout?
 - Go-Back-N: resend all unacknowledged data on timeout
 - Selective Repeat: timer per packet, resend as needed

54

SLIDING WINDOW: RECEIVER



- Receiver buffers too:
 - data may arrive out-of-order
 - or faster than can be consumed
 - Flow control: tell sender how much buffer left at receiver
- Receiver ACK choices:
 - Cumulative, Selective (exempt missing frames), Negative

55

DECIDING WHEN TO RETRANSMIT

- How do you know when a packet has been lost?
 - Ultimately sender uses timers to decide when to retransmit
- But how long should the timer be?
 - Too long: inefficient (large delays, poor use of bandwidth)
 - Too short: may retransmit unnecessarily (causing extra traffic)
- Right timer is based on the round-trip time (RTT)
 - Which can vary greatly

56

